

Laborator 6

```
biancapinghireac@vbox:~/S0/lab6/src$ make lib
gcc -Wall -g -O -c -o error.o error.c
gcc -Wall -g -O -c -o prexit.o prexit.c
gcc -Wall -g -O -c -o tellwait.o tellwait.c
ar rcs liblab6.a error.o prexit.o tellwait.o
biancapinghireac@vbox:~/S0/lab6/src$ make
gcc -o fork1 fork1.c liblab6.a
gcc -o wait1 wait1.c liblab6.a
gcc -o fork2 fork2.c liblab6.a
gcc -o tellwait1 tellwait1.c liblab6.a
gcc -o tellwait2 tellwait2.c liblab6.a
gcc -o echoall echoall.c liblab6.a
gcc -o exec1 exec1.c liblab6.a
```

Crearea bibliotecii.

```
biancapinghireac@vbox:~/S0/lab6/src$ ./fork1
a write to stdout
before fork
pid = 4454, glob = 7, var = 89
pid = 4453, glob = 6, var = 88
biancapinghireac@vbox:~/S0/lab6/src$ ./fork1 > temp.out
biancapinghireac@vbox:~/S0/lab6/src$ cat temp.out
a write to stdout
before fork
pid = 4472, glob = 7, var = 89
before fork
pid = 4471, glob = 6, var = 88
```

```
#include <sys/types.h>
#include "ourhdr.h"

int glob = 6; /* external variable in initialized data */
char buf[] = "a write to stdout\n";

int main(void)
{
    int var; /* automatic variable on the stack */
    pid_t pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) { /* child */
        glob++; /* modify variables */
        var++;
    } else
        sleep(2); /* parent */

    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```

Fork1.c

Cu `write()` scrie direct in `stdout`, pentru ca nu da flush la buffer, cand apeleaza `printf()`, in terminal se va afisa o singura data(se da flush automat), in schimb in fisier nu se da flush buffer-ului si se compiaza atat in procesul parinte cat si in cel fiu prin `fork()`.

`Sleep()` pentru parinte lasa timp procesului copil sa actioneze(increment-ul valorilor) fara a influenta si valorile din cel parinte.

```
biancapinghireac@vbox:~/S0/lab6/src$ ./wait1
normal termination, exit status = 7
abnormal termination, signal number = 6 (core file generated)
wait error: No child processes
normal termination, exit status = 1
```

```
#include <sys/types.h>
#include <sys/wait.h>
#include "ourhdr.h"
#include "prexit.h"

int main(void)
{
    pid_t    pid;
    int      status;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0)          /* child */
        exit(7);

    if (wait(&status) != pid)  /* wait for child */
        err_sys("wait error");
    pr_exit(status);           /* and print its status */

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0)          /* child */
        abort();               /* generates SIGABRT */

    if (wait(&status) != pid)  /* wait for child */
        err_sys("wait error");
    pr_exit(status);           /* and print its status */

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0)          /* child */
        status = 0;           /* divide by 0 generates SIGFPE */

    if (wait(&status) != pid)  /* wait for child */
        err_sys("wait error");
    pr_exit(status);           /* and print its status */

    exit(0);
}
```

Wait1.c

Primul copil se termina normal cu **exit(7);**.

Al doilea copil se termina anormal prin semnal (**abort()** → **SIGABRT**).

Al treilea copil ar trebui sa produca o exceptie de tip diviziune la zero (**SIGFPE**), in acest caz codul zice doar **status = 0;** .

```
biancapinghireac@vbox:~/S0/lab6/src$ ./fork2
biancapinghireac@vbox:~/S0/lab6/src$ second child, parent pid = 1920
```

```
#include <sys/types.h>
#include <sys/wait.h>
#include "ourhdr.h"

int main(void)
{
    pid_t    pid;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {          /* first child */
        if ( (pid = fork()) < 0)
            err_sys("fork error");
        else if (pid > 0)
            exit(0);             /* parent from second fork == first child */

        /* We're the second child; our parent becomes init as soon
           as our real parent calls exit() in the statement above.
           Here's where we'd continue executing, knowing that when
           we're done, init will reap our status. */

        sleep(2);
        printf("second child, parent pid = %d\n", getppid());
        exit(0);
    }

    if (waitpid(pid, NULL, 0) != pid) /* wait for first child */
        err_sys("waitpid error");

    /* We're the parent (the original process); we continue executing,
       knowing that we're not the parent of the second child. */
    exit(0);
}
```

Fork2.c

In exemplul de mai sus se creaza un proces orfan(second child) care va fi preluat de systemd/init

```
biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait1
ouuuttpuutt ffrroomm pcahrielndt

biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait1
output from parent
output from child

biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait1
output from parent
output from child

biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait1
output from child
output from parent
```

```
#include <sys/types.h>
#include "ourhdr.h"

static void charatime(char *);

int main(void)
{
    pid_t  pid;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
    }
    exit(0);
}

static void charatime(char *str)
{
    char    *ptr;
    int     c;

    setbuf(stdout, NULL);           /* set unbuffered */
    for (ptr = str; c = *ptr++; )
        putc(c, stdout);
}
```

Tellwait1.c

In exemplul de mai sus se observa lipsa sincronizarii proceselor care incearca sa modifice aceeaasi zona de memorie (buffer ul stdout). Din acest motiv trebuie sincronizate.

```
biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait2
output from parent
output from child
```

```

#include <sys/types.h>
#include "ourhdr.h"

static void charatime(char *);

int main(void)
{
    pid_t    pid;

    TELL_WAIT();

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {
        WAIT_PARENT();          /* parent goes first */
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
        TELL_CHILD(pid);
    }
    exit(0);
}

static void charatime(char *str)
{
    char      *ptr;
    int       c;

    setbuf(stdout, NULL);      /* set unbuffered */
    for (ptr = str; c = *ptr++; )
        putc(c, stdout);
}

```

Tellwait2.c

Aici se observa sincronizarea, deoarece copilul asteapta ca parintele sa termine procesul, apoi incepe si el.

```

biancapinghireac@vbox:~/S0/lab6/src$ ./exec1
argv[0]: echoall
argv[1]: myarg1
argv[2]: MY ARG2
USER=unknown
PATH=/tmp
exec1p error: No such file or directory

```

```

#include <sys/types.h>
#include <sys/wait.h>
#include "ourhdr.h"

char *env_init[] = { "USER=unknown", "PATH=/tmp", NULL };

int main(void)
{
    pid_t pid;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) { /* specify pathname, specify environment */
        if (execl("./echoall",
                  "echoall", "myarg1", "MY ARG2", (char *) 0,
                  env_init) < 0)
            err_sys("execl error");
    }
    if (waitpid(pid, NULL, 0) < 0)
        err_sys("wait error");

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) { /* specify filename, inherit environment */
        if (execlp("echoall",
                  "echoall", "only 1 arg", (char *) 0) < 0)
            err_sys("execlp error");
    }
    exit(0);
}

```

Exec1.c

Programul creeaza doi copii cu fork. Primul foloseste **execl** pentru a rula **./echoall** cu argumente si un mediu (environment) propriu, unde PATH este doar **/tmp**. Functioneaza deoarece se da calea explicita spre executabil.

Al doilea copil foloseste **execlp** cu numele programului, dar nu il gaseste, pentru ca in PATH exista doar **/tmp**, iar acolo nu e **echoall**. De aceea apare eroare. Exemplul arata diferenta dintre **execl** (cale si environment explicit) si **execlp** (moșteneste environment si cauta in PATH).

fork(2) System Calls Manual [fork\(2\)](#)

NAME fork - create a child process

LIBRARY standard C library ([libc](#), [-lc](#))

SYNOPSIS

```
#include <unistd.h>
pid_t fork(void);
```

DESCRIPTION

fork() creates a new process by duplicating the calling process. The new process is referred to as the **child** process. The calling process is referred to as the **parent** process.

The child process and the parent process run in separate memory spaces. At the time of **fork()** both memory spaces have the same content. Memory writes, file mappings ([mmap\(2\)](#)), and unmappings ([munmap\(2\)](#)) performed by one of the processes do not affect the other.

The child process is an exact duplicate of the parent process except for the following points:

- The child has its own unique process ID, and this PID does not match the ID of any existing process group ([setpgid\(2\)](#)) or session.
- The child's parent process ID is the same as the parent's process ID.
- The child does not inherit its parent's memory locks ([mlock\(2\)](#), [mlockall\(2\)](#)).
- Process resource utilizations ([getrusage\(2\)](#)) and CPU time counters ([times\(2\)](#)) are reset to zero in the child.
- The child's set of pending signals is initially empty ([sigpending\(2\)](#)).
- The child does not inherit semaphore adjustments from its parent ([semop\(2\)](#)).
- The child does not inherit process-associated record locks from its parent ([fcntl\(2\)](#)). (On the other hand, it does inherit [fcntl\(2\)](#) open file description locks and [flock\(2\)](#) locks from its parent.)
- The child does not inherit timers from its parent ([setitimer\(2\)](#), [alarm\(2\)](#), [timer_create\(2\)](#)).
- The child does not inherit outstanding asynchronous I/O operations from its parent ([aio_read\(3\)](#), [aio_write\(3\)](#)), nor does it

Manual page: fork(2). Line 1. (press h for help or q to quit)

Fork

```

wait(2)                                     System Calls Manual                                     wait(2)

NAME
    wait, waitpid, waitid - wait for process to change state

LIBRARY
    Standard C library (libc, libc)

SYNOPSIS
#include <sys/wait.h>

pid_t wait(int *Wunltable wstatus);
pid_t waitpid(pid_t pid, int *Wunltable wstatus, int options);
int waitid(idtype_t idtype, id_t id, siginfo_t *Wunltable info, int options);
/* This is the glibc and POSIX interface; see
   NOTES for information on the raw system call. */

Feature Test Macro Requirements for glibc (see feature_test_macros(7)):

wait(2):
    Since glibc 2.26:
        _XOPEN_SOURCE >= 500 || _POSIX_C_SOURCE >= 200809L
    glibc 2.25 and earlier:
        _XOPEN_SOURCE
        || /* Since glibc 2.12: */ _POSIX_C_SOURCE >= 200809L
        || /* glibc <= 2.19: */ _BSD_SOURCE

DESCRIPTION
    All of these system calls are used to wait for state changes in a child of the calling process, and obtain information about the child whose state has changed. A state change is considered to be: the child terminated; the child was stopped by a signal; or the child was resumed by a signal. In the case of a terminated child, performing a wait allows the system to release the resources associated with the child; if a wait is not performed, then the terminated child remains in a "zombie" state (see NOTES below).

    If a child has already changed state, then these calls return immediately. Otherwise, they block until either a child changes state or a signal handler interrupts the call (assuming that system calls are not automatically restarted using the SA_RESTART flag of sigaction(2)). In the remainder of this page, a child whose state has changed and which has not yet been waited upon by one of these system calls is termed waitable.

wait() and waitpid()
    The wait() system call suspends execution of the calling thread until one of its children terminates. The call wait(2):
Manual page wait(2) line 1 (press h for help or q to quit)

```

Wait

```

_exit(2)                                     System Calls Manual                                     _exit(2)

NAME
    _exit, _Exit - terminate the calling process

LIBRARY
    Standard C library (libc, libc)

SYNOPSIS
#include <unistd.h>

[[noreturn]] void _exit(int status);
#include <stdlib.h>
[[noreturn]] void _Exit(int status);

Feature Test Macro Requirements for glibc (see feature_test_macros(7)):

_exit(2):
    _ISOC99_SOURCE || _POSIX_C_SOURCE >= 200112L

DESCRIPTION
    _exit() terminates the calling process "immediately". Any open file descriptors belonging to the process are closed. Any children of the process are inherited by init(1) (or by the nearest "subreaper" process as defined through the use of the prctl(2) PR_SET_CHILD_SUBREAPER operation). The process's parent is sent a SIGCHLD signal.

    The value status & 0xFF is returned to the parent process as the process's exit status, and can be collected by the parent using one of the wait(2) family of calls.

    The function _Exit() is equivalent to _exit().

RETURN VALUE
    These functions do not return.

STANDARDS
    _exit()
        POSIX.1-2008.

    _Exit()
        C11, POSIX.1-2008.

Manual page _exit(2) line 1 (press h for help or q to quit)

```

Exit

EXERCITIUL 2:

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/wait.h>

#define BUFFER_SIZE 1024

void main(){
    char command[BUFFER_SIZE];
    int status;
    while(1){
        printf("$ ");
        fgets(command, BUFFER_SIZE, stdin);
        command[strlen(command)-1]='\0';

        if(strcmp(command,"exit")==0)
            break;

        pid_t pid = fork(); //proces nou
        if(pid==-1){ //procesul nu s-a creat corespunzator
            exit(20); //20 fork
        }else if(pid ==0){
            if(execlp(command,command,(char *) NULL) ==-1){
                exit(21); //21 command error
            }
        }
        else{
            waitpid(pid,&status,0);
        }
    }
    exit(1);
}

```

Citeste comenzile de la tastatura si le ruleaza intr-un proces copil folosind `fork` si `execlp`.
 Daca scrii `exit`, se opreste. Daca sunt erori la fork sau la exec, programul iese cu un cod de eroare special.te


```

biancapinghireac@vbox:~/S0/lab6/src$ ./shell
$ tree
.
├── echoall
├── echoall.c
├── error.c
├── error.o
├── exec1
├── exec1.c
├── fork1
├── fork1.c
├── fork2
├── fork2.c
├── liblab6.a
├── Makefile
├── ourhdr.h
├── prexit.c
├── prexit.h
├── prexit.o
├── shell
├── shell.c
├── tellwait1
├── tellwait1.c
├── tellwait2.c
├── tellwait.c
├── tellwait.o
├── temp.out
├── wait1
└── wait1.c

1 directory, 26 files
$

```

EXERCITIUL 3:

```

biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait2
output from parent
output from child

```

Nemodificat

```

biancapinghireac@vbox:~/S0/lab6/src$ gcc tellwait2.c -o tellwait2
biancapinghireac@vbox:~/S0/lab6/src$ ./tellwait2
output from child
output from parent

```

Modificat

Cod:

```

GNU nano 8.1 tellwait2.c
#include <sys/types.h>
#include "ourhdr.h"
#include <stdio.h>

static void charatime(char *);

int main(void)
{
    pid_t  pid;

    TELL_WAIT();

    if ( (pid = fork()) < 0)
        perror("fork error");
    else if (pid == 0) {
        charatime("output from child\n");
        TELL_PARENT(pid);
    } else {
        WAIT_CHILD();          /* child goes first */
        charatime("output from parent\n");
    }
    exit(0);
}

static void charatime(char *str)
{
    char    *ptr;
    int     c;

    setbuf(stdout, NULL);      /* set unbuffered */
    for (ptr = str; c = *ptr++; )
        putc(c, stdout);
}

```

Am inversat ordinea proceselor(acum procesul parinte asteapta terminarea procesului fiu).